



US009342256B2

(12) **United States Patent**
Peterson et al.

(10) **Patent No.:** **US 9,342,256 B2**
(45) **Date of Patent:** **May 17, 2016**

(54) **EPOCH BASED STORAGE MANAGEMENT
FOR A STORAGE DEVICE**

(56) **References Cited**

U.S. PATENT DOCUMENTS

(71) Applicant: **Fusion-io, Inc.**, Salt Lake City, UT (US)

(72) Inventors: **James G. Peterson**, San Jose, CA (US);
Nisha Talagala, Livermore, CA (US);
Swaminathan Sundararaman, San
Jose, CA (US); **Sriram Subramanian**,
San Jose, CA (US)

(73) Assignee: **SANDISK TECHNOLOGIES, INC.**,
Plano, TX (US)

(*) Notice: Subject to any disclaimer, the term of this
patent is extended or adjusted under 35
U.S.C. 154(b) by 370 days.

(21) Appl. No.: **13/831,448**

(22) Filed: **Mar. 14, 2013**

(65) **Prior Publication Data**

US 2014/0281307 A1 Sep. 18, 2014

(51) **Int. Cl.**
G06F 3/06 (2006.01)
G06F 17/30 (2006.01)
G06F 11/14 (2006.01)

(52) **U.S. Cl.**
CPC **G06F 3/065** (2013.01); **G06F 3/0619**
(2013.01); **G06F 3/0688** (2013.01); **G06F**
11/1402 (2013.01); **G06F 17/30088** (2013.01)

(58) **Field of Classification Search**

None

See application file for complete search history.

* cited by examiner

Primary Examiner — Reginald Bragdon

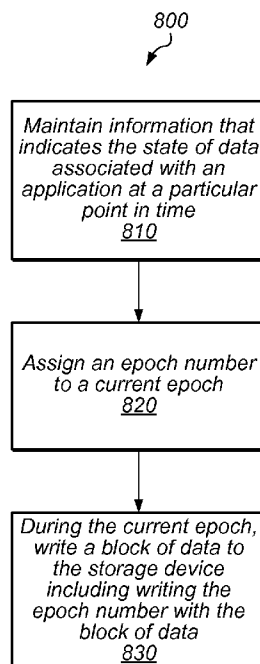
Assistant Examiner — Edward Wang

(74) *Attorney, Agent, or Firm* — Kunzler Law Group, PC

(57) **ABSTRACT**

Techniques are disclosed relating to handling snapshot data for a storage device. In one embodiment, a computing system maintains information that indicates the state of data associated with an application at a particular point in time. In this embodiment, the computing system assigns an epoch number to a current epoch, where the current epoch is an interval between the particular point in time and a future point in time. In this embodiment, the computing system writes, during the current epoch, a block of data to the storage device. In this embodiment, the writing the block of data includes storing the epoch number with the block of data.

21 Claims, 13 Drawing Sheets



Computing
System
100

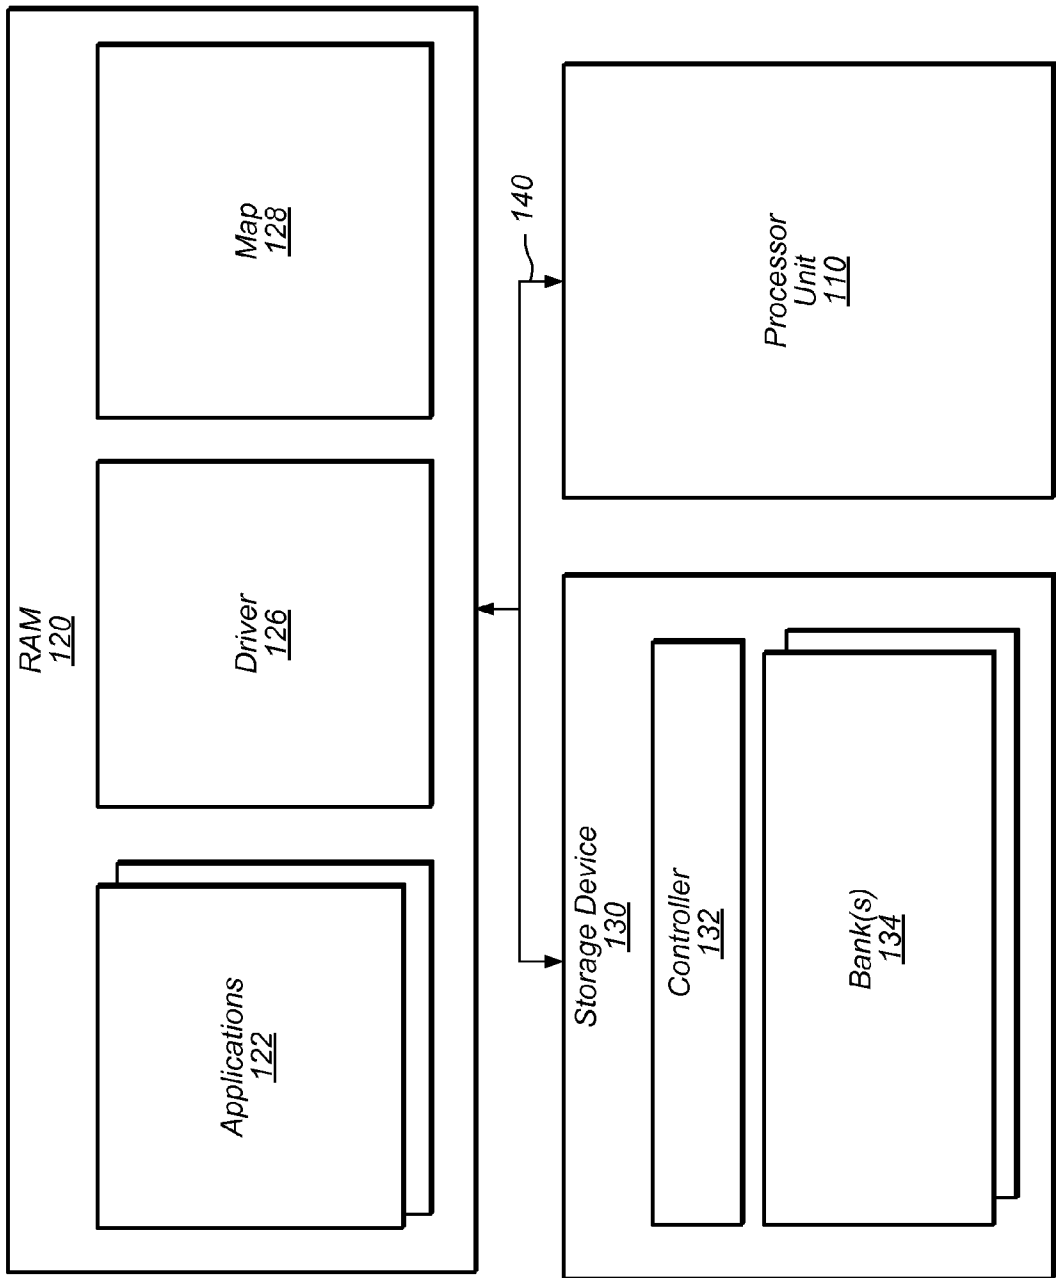


FIG. 1




FIG. 2A





FIG. 2B

Logical Address Space 302



Physical Address Space 304

FIG. 3A




FIG. 3B




FIG. 3C




FIG. 4




FIG. 5

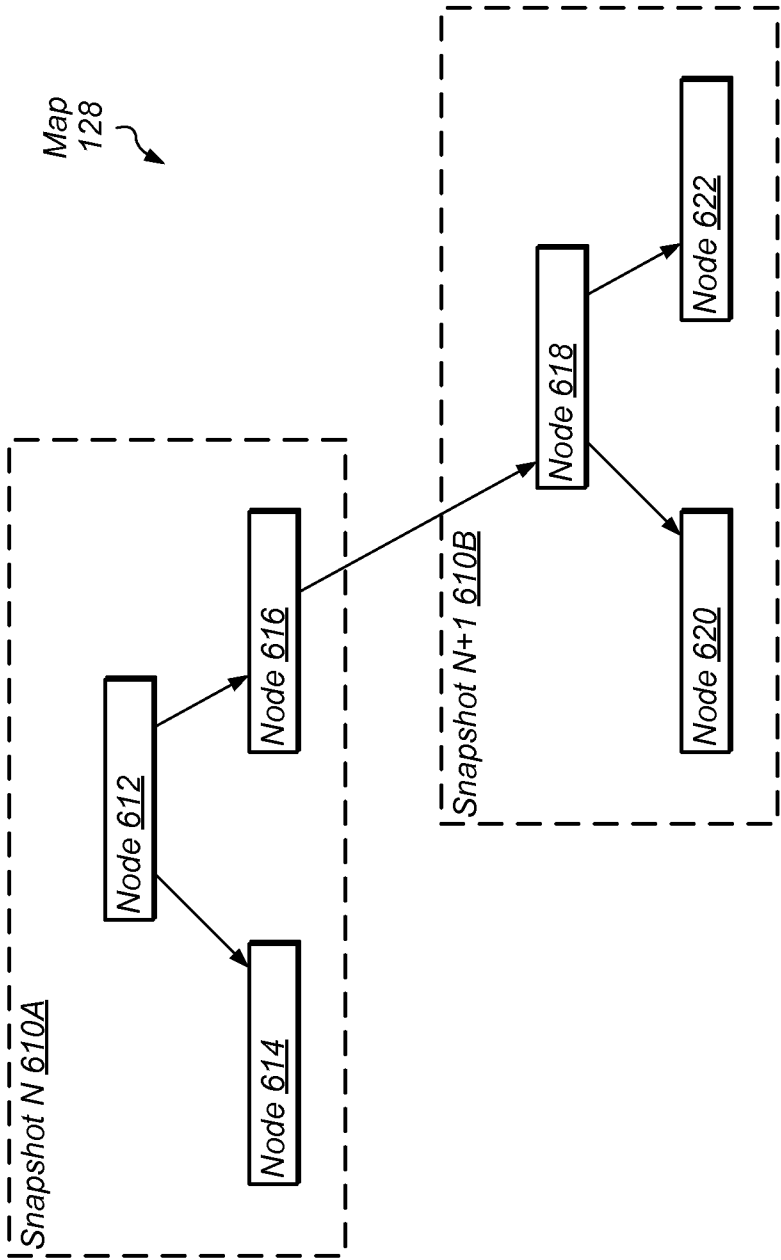


FIG. 6




FIG. 7

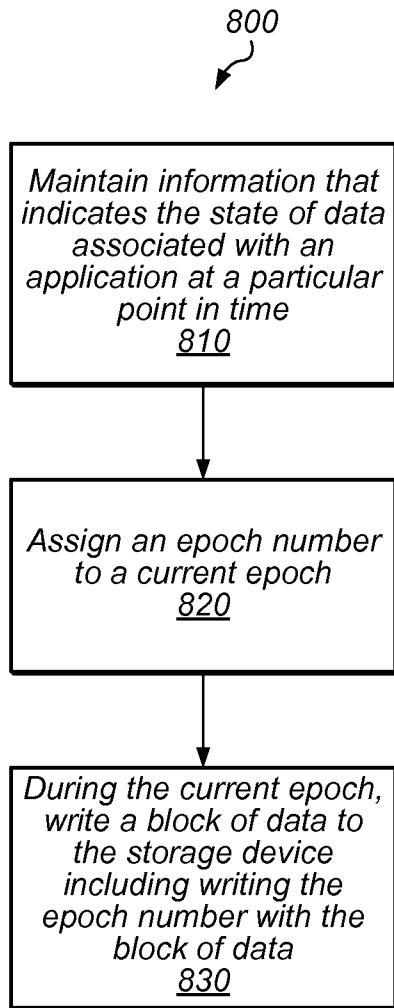


FIG. 8A

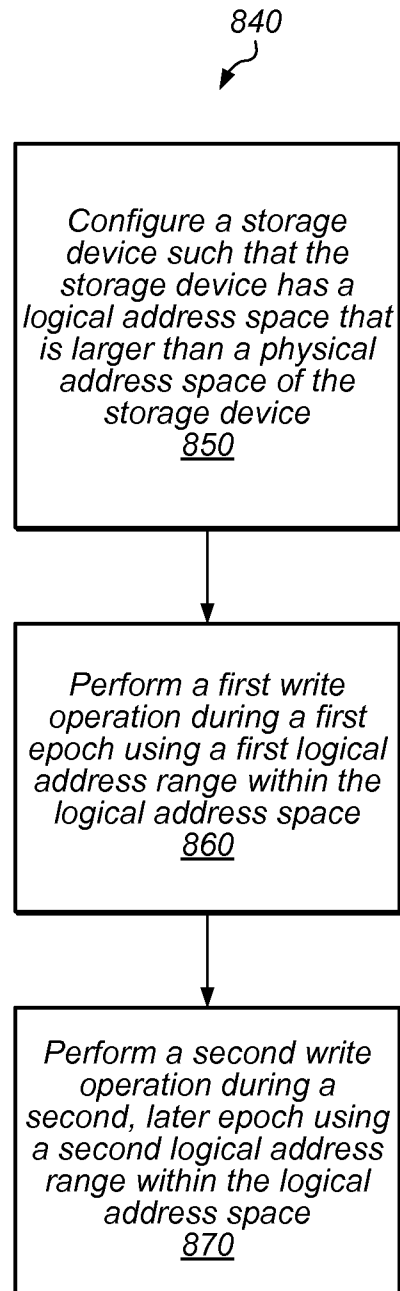


FIG. 8B




FIG. 9

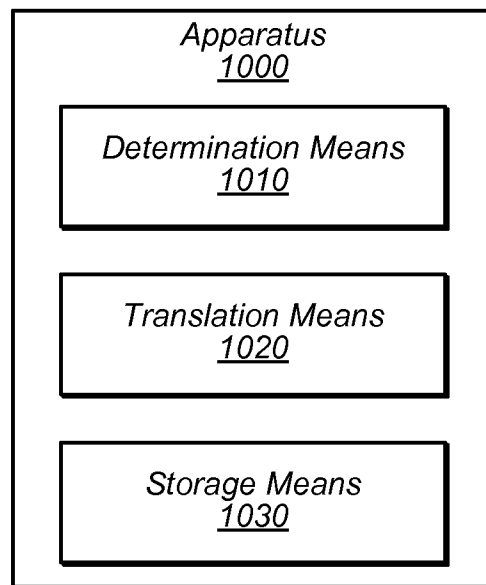


FIG. 10A

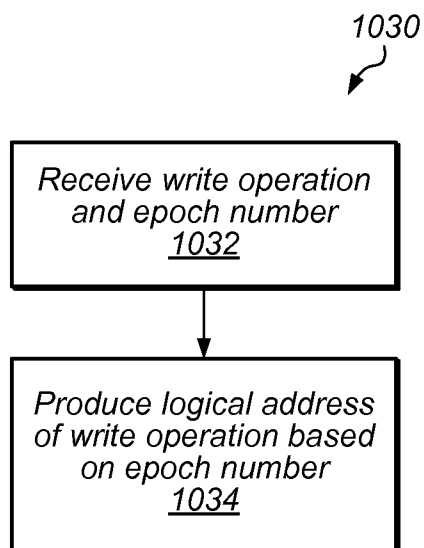


FIG. 10B

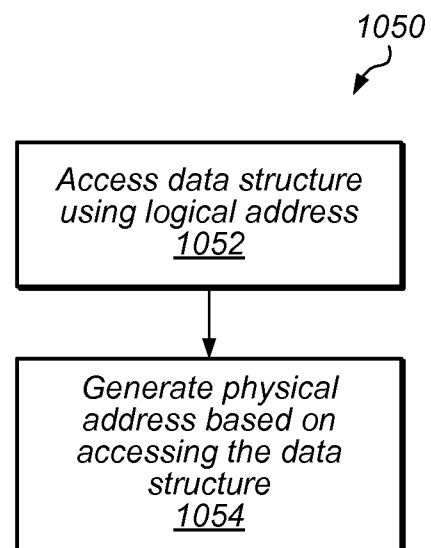


FIG. 10C

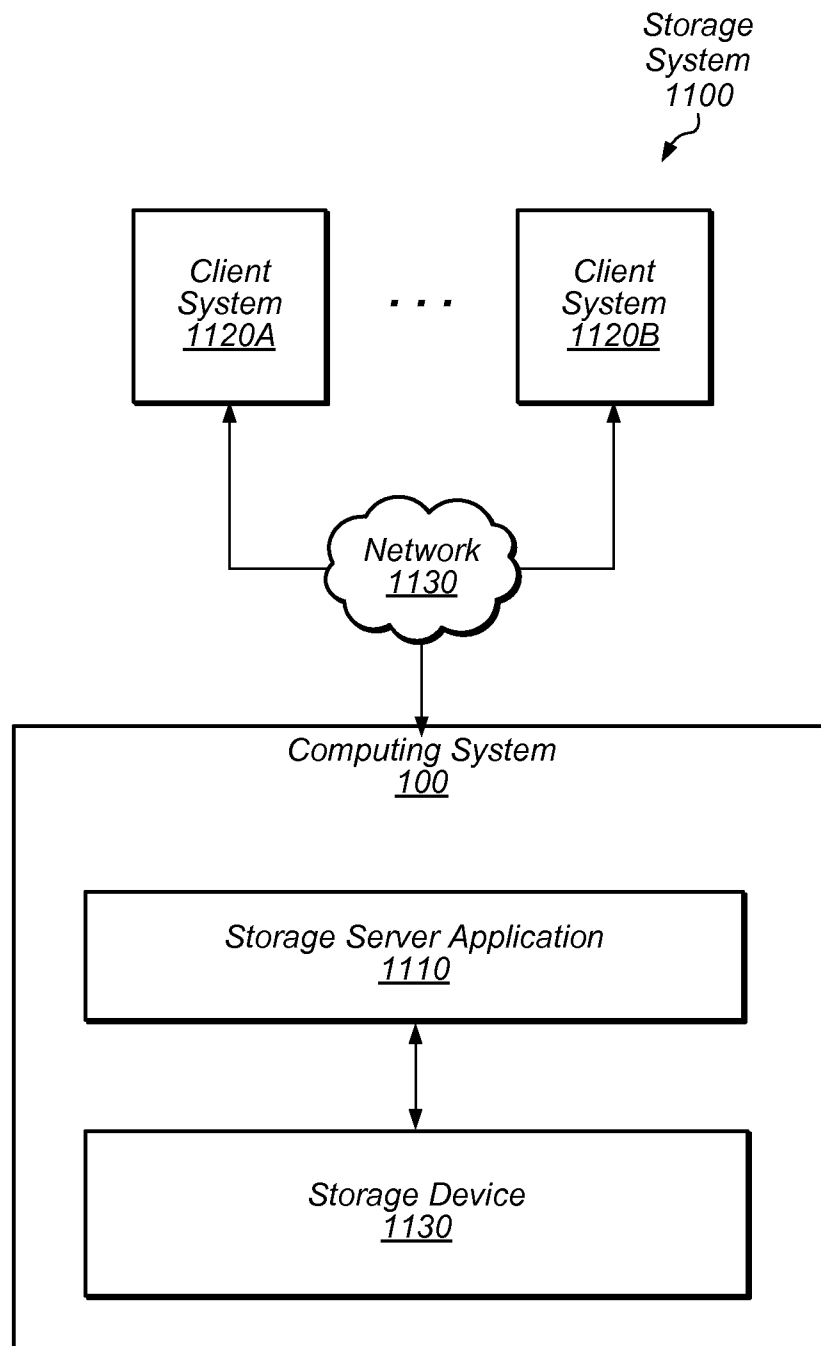


FIG. 11

1

EPOCH BASED STORAGE MANAGEMENT FOR A STORAGE DEVICE

BACKGROUND

1. Technical Field

This disclosure relates generally to accessing data on a physical recording medium, and more specifically to handling snapshot information on such a medium.

2. Description of the Related Art

Modern storage systems often require various restore capabilities for stored data. Such a restoration may be performed, for example, to facilitate crash recovery in the event of power loss, to recover a last known valid state in the event of data corruption, etc. This functionality may be achieved by taking “snapshots” of data for particular applications such that a snapshot includes information indicating the current state of data for the application at the time the snapshot is taken. For example, taking the snapshot of a database application might include copying the current values of data within the database to an alternate storage.

In the context of flash-based storage, to improve the longevity of memory cells, modern storage systems may implement a log-structured storage to ensure that writes to cells are more evenly distributed across the storage to produce better wear leveling (as opposed to writing particular cells frequently while other cells go unused). When storing data using a log-structure storage, data may be written at an append point that starts at an initial portion in the storage and advances forward as writes are performed. A driver may map logical addresses used by an application to physical locations on a storage device where the data is actually stored, e.g., according to a log-structured implementation.

SUMMARY

The present disclosure describes embodiments of techniques for handling snapshot data for a storage device. These techniques may allow for restoration of data and snapshot data when data corruption occurs. These techniques may allow for efficient storage of snapshot data by storing a delta of information written during an epoch associated with each snapshot.

In one embodiment, a computing system maintains information that indicates the state of data associated with an application at a particular point in time. In this embodiment, the computing system assigns an epoch number to a current epoch, where the current epoch is an interval between the particular point in time and a future point in time. In this embodiment, the computing system writes, during the current epoch, a block of data to the storage device. In this embodiment, the writing the block of data includes storing the epoch number with the block of data.

In another embodiment, an apparatus is disclosed that includes an association module and a storage module. In this embodiment, the association module is configured to associate a logical address range for a storage device with an epoch number. In this embodiment, the epoch number is associated with write operations to the storage device during a particular epoch. In this embodiment, the storage module is configured to handle one or more storage operations associated with the epoch number using the logical address range. For example, write during an epoch may be associated with the epoch number of the epoch. Read operations may be associated with the current epoch or a particular snapshot.

In yet another embodiment, a non-transitory computer readable medium has program instructions stored thereon.

2

The program instructions are executable by a computing system to cause the computing system to perform operations. The operations include configuring a storage device such that the storage device has a logical address space that is larger than a physical address space of the storage device. The operations further include performing a first write operation during a first epoch using a first logical address range within the logical address space and performing a second write operation during a second, later epoch using a second logical address range within the logical address space.

In yet another embodiment, an apparatus is disclosed that includes first, second, and third means. The first means is for determining a logical address for a write operation based on a current epoch number. The second means is for translating the logical address to a physical address on a storage device. The third means is for storing the current epoch number on the storage device.

BRIEF DESCRIPTION OF THE DRAWINGS

FIG. 1 is a block diagram illustrating one embodiment of a computing system that supports storage snapshots.

FIG. 2A is a diagram illustrating exemplary snapshot data.

FIG. 2B is a diagram illustrating an exemplary data structure for maintaining relationships between snapshots.

FIGS. 3A-3C are block diagrams illustrating embodiments of logical and physical address spaces.

FIG. 4 is a block diagram illustrating one embodiment of a driver for the storage device.

FIG. 5 is a block diagram illustrating one embodiment of a map data structure.

FIG. 6 is a block diagram illustrating another embodiment of a map data structure.

FIG. 7 is a block diagram illustrating another embodiment of a driver for the storage device.

FIGS. 8A and 8B are flow diagrams illustrating embodiments of methods.

FIG. 9 is a block diagram illustrating one embodiment of an apparatus having an association module and a storage module.

FIG. 10A is a block diagram illustrating another embodiment of an apparatus having a determination means, translation means, and storage means.

FIGS. 10B and 10C are flow diagrams illustrating embodiments of algorithms implemented by a determination means and a translation means.

FIG. 11 is a block diagram illustrating one embodiment of a storage system.

This specification includes references to “one embodiment” or “an embodiment.” The appearances of the phrases “in one embodiment” or “in an embodiment” do not necessarily refer to the same embodiment. Particular features, structures, or characteristics may be combined in any suitable manner consistent with this disclosure.

This disclosure also includes and references the accompanying drawings. In the drawings, similar symbols typically identify similar components, unless context dictates otherwise. The illustrative embodiments described in the detailed description, drawings, and claims are not meant to be limiting. Other embodiments may be utilized, and other changes may be made to these exemplary embodiments, without departing from the scope of the disclosure.

Various units, circuits, or other components in this disclosure may be described or claimed as “configured to” perform a task or tasks. In such contexts, “configured to” is used to connote structure by indicating that the units/circuits/components include structure (e.g., circuitry) that performs those

task or tasks during operation. As such, the unit/circuit/component can be said to be configured to perform the task even when the specified unit/circuit/component is not currently operational (e.g., is not on). The units/circuits/components used with the “configured to” language include hardware—for example, circuits, memory storing program instructions executable to implement the operation, etc. Additionally, “configured to” can include generic structure (e.g., generic circuitry) that is manipulated by software and/or firmware (e.g., an FPGA or a general-purpose processor executing software) to operate in a manner that is capable of performing the task(s) at issue. Reciting that a unit/circuit/component is “configured to” perform one or more tasks is expressly intended not to invoke 35 U.S.C. §112, sixth paragraph, for that unit/circuit/component.

As used herein, the term “based on” is used to describe one or more factors that affect a determination. This term does not foreclose additional factors that may affect a determination. That is, a determination may be solely based on those factors or based, at least in part, on those factors. Consider the phrase “determine A based on B.” While in this case, B is a factor that affects the determination of A, such a phrase does not foreclose the determination of A from also being based on C. In other instances, A may be determined based solely on B.

DETAILED DESCRIPTION

The disclosure initially describes, with reference to FIGS. 1 and 2A-2B, a computing system that handles snapshot data for a storage system. To facilitate this description, logical and physical address spaces associated with the storage device are described with reference to FIGS. 3A-3C. Drivers and map structures usable to handle snapshot data within the storage device are described with reference to FIGS. 4-7. Further embodiments of techniques for handling snapshot data are described in further detail with reference to FIGS. 8-11.

Turning now to FIG. 1, a block diagram of a computing system 100 that supports snapshots is depicted. Computing system 100 may be any suitable type of computing device such as a server, laptop, desktop, a mobile device, etc. In some embodiments, computing system 100 may include multiple computing devices working together. For example, in one embodiment, computing system 100 may be multiple servers coupled together at a data center configured to store data on behalf of multiple clients, such as the storage system discussed below in conjunction with FIG. 11. In the illustrated embodiment, computing system 100 includes a processor unit 110, random access memory (RAM) 120, and storage device 130 coupled together via an interconnect 140. As shown, RAM 120 may include program instructions for one or more applications 122. RAM may include a driver 126 for storage device 130, which, in turn, may include a controller 132 and one or more storage banks 134. RAM 120 may also include a map 128.

In various embodiments, driver 126 is described as having various functionality. This functionality may be implemented in software, hardware or a combination thereof. Such functionality may be implemented within an application 122, in one embodiment. In another embodiment, this functionality may be implemented by software stored within a memory of controller 132 and executed by a processor of controller 132. In still another embodiment, controller 132 may include dedicated circuitry to implement functionality of driver 126. In sum, the depiction of driver 126 as being implemented in RAM 120 should not be seen as limiting, but rather as a

depiction of an exemplary embodiment. Similarly, map 128 may be stored in driver 126, controller 132, and/or any other appropriate location.

Storage device 130 is representative of any physical medium upon which data can be recorded. As used herein, the term “recorded” refers broadly to the process of an electronic computing device storing, writing or otherwise transferring one or more data values on to some physical recording medium for subsequent use. Accordingly, a “physical recording medium” is used herein to refer to any medium on which data may be recorded by an electronic computing device. Further, the terms “storage” and “memory” are used herein to be synonymous with “physical recording medium.” Given this broad definition, the designations memory (when referring to RAM 120) and storage (when referring to storage device 130) in FIG. 1 and elsewhere in this disclosure may refer to volatile and/or non-volatile media. Such media may also be referred to herein as “memory,” and portions of such media may be referred to as “blocks,” “cells,” “storage blocks,” “memory blocks,” etc. Collectively, a group of these blocks may be referred to as a “storage array,” “memory array,” etc.

References in this disclosure to “accessing” data in storage device 130 or “storage requests” to storage device 130 refer to any type of transaction, including writing data to storage device 130 and/or reading data from storage device 130, as well as, TRIM operations, maintenance accesses, discovery accesses, load and store operations under memory semantics, and the like. Further, given the broad definitions of “storage” and “memory” referred to above, these accesses may be applicable to a storage device that has non-volatile memory and/or volatile components.

In some embodiments, storage device 130 may be implemented such that it includes non-volatile memory. Accordingly, in such an embodiment, storage banks 134 may include non-volatile storage devices such as hard disk drives (e.g., Integrated Drive Electronics (IDE) drives, Small Computer System Interface (SCSI) drives, Serial Attached SCSI (SAS) drives, Serial AT Attachment (SATA) drives, etc.), tape drives, writable optical drives (e.g., CD drives, DVD drives, Blu-Ray drives, etc.) etc.

In some embodiments, storage device 130 may be implemented such that it includes non-volatile solid-state memory. Accordingly, in such an embodiment, storage banks 134 may include any suitable type of solid-state storage media including, but not limited to, NAND flash memory, NOR flash memory, nano RAM (“NRAM”), magneto-resistive RAM (“MRAM”), phase change RAM (“PRAM”), Racetrack memory, Memristor memory, nanocrystal wire-based memory, silicon-oxide based sub-10 nanometer process memory, graphene memory, Silicon-Oxide-Nitride-Oxide-Silicon (“SONOS”), Resistive random-access memory (“RRAM”), programmable metallization cell (“PMC”), conductive-bridging RAM (“CBRAM”), etc. In some embodiments, storage banks 134 may include multiple, different types of solid-state storage media.

In other embodiments, storage device 130 may be implemented such that it includes volatile memory. Storage banks 134 may thus correspond to any suitable volatile memory including, but not limited to such as RAM, dynamic RAM (DRAM), static RAM (SRAM), synchronous dynamic RAM (SDRAM), etc. Although shown independently of processor unit 110, in some embodiments, storage device 130 may correspond to memory within processor unit 110 such as one or more cache levels (e.g., L1, L2, L3, etc.) within processor unit 110.

In sum, various functionality will be described herein pertaining to storage device **130**. Such functionality may be applicable to any suitable form of memory including both non-volatile and volatile forms. Thus, while particular embodiments of driver **126** are described herein within the context of non-volatile solid-state memory arrays, driver **126** may also be applicable to other recording media such as volatile memories and other types of non-volatile memories, particularly those that include a reclamation process.

Controller **132**, in one embodiment, is configured to manage operation of storage device **130**. Accordingly, controller **132** may facilitate performance of read operations at specified addresses (e.g., “physical addresses” as discussed below) including selecting the appropriate banks **134** and accessing the data within the appropriate cells within those banks. Controller **132** may facilitate performance of write operations including programming of particular cells. Controller **132** may also perform preparation operations to permit subsequent writes to storage device **130** such as, in one embodiment, erasing blocks of cells for subsequent reuse. (The cycle of programming and erasing a block of cells may be referred to as a “PE cycle.”) In some embodiments, controller **132** implements separate read and write data pipelines to perform read and write operations in parallel. In one embodiment, controller **132** is also configured to communicate with driver **126** (discussed below) over interconnect **140**. For example, in some embodiments, controller **132** communicates information for read and write operations via direct memory access (DMA) transactions coordinated by a DMA controller. Accordingly, controller **132** may support any suitable interconnect type such as a peripheral component interconnect (PCI), PCI express (PCI-e), serial advanced technology attachment (“serial ATA” or “SATA”), parallel ATA (“PATA”), small computer system interface (“SCSI”), IEEE 1394 (“FireWire”), Fiber Channel, universal serial bus (“USB”), etc. In some embodiments, controller **132** may also perform other operations such as error checking, data compression, encryption and decryption, packet assembly and disassembly, etc.

In various embodiments, storage device **130** is organized as a log-structured storage. As used herein, the term “log structure” refers to an arrangement of data on a storage medium in which an append point is used to determine where data is stored; the append point is advanced sequentially through an “address space” as data is stored. A log-structured storage is simply a storage device that is organized using a log structure. The use of a log structure also connotes that metadata is stored in conjunction with the data in order to permit the storage device **130** to be restored to a previous state. Such a restoration may be performed, for example, to facilitate crash recovery in the event of power loss, to recover a last known valid state in the event of data corruption, etc. As used herein, the term “address space” refers to a range of addresses that can be used to specify data within a storage device. As will be described below, a log-structured storage may have both logical and physical address spaces. The term “logical address space” refers to an address space as perceived by higher-level processes even though this address space may not be representative of how data is actually organized on the physical media of storage device **130** or the actual number of physical address locations actually in use, reserved, or allocated to a higher-level process. In contrast, the term “physical address space” refers to the address space used by lower-level processes and may be indicative of how data is organized on the physical media of storage device **130** and the actual number of physical address locations in use by a higher-level process. Embodiments of logical and physical address spaces are dis-

cussed in further detail in conjunction with FIGS. **3A** and **3B**, respectively. One embodiment of a log structure is discussed in conjunction with FIG. **3C**.

In various embodiments, using a log structure may permit multiple instances of a set of data to be present in storage device **130** as the data is written, modified, and rewritten to storage. As part of tracking data in a physical address space, older instances of stored data (i.e., those instances that are not the current instance) may be indicated as invalid. For example, in one embodiment, when a value is to be updated, the value may be written at a storage block specified by the current append point (rather than at the location where the value was previously stored). In response to the write being successfully performed, any previously stored instances of that value may be marked as invalid. As used herein, the term “invalid” refers to data that no longer needs to be stored by the system (e.g., because a newer copy of the data exists). Similarly, the term “invalidating” refers to the marking of data as invalid (e.g., storing a record in a data structure) or to the storing of an instance of data when a previous instance of the data existed in storage device **130**, the storing making the previous instance invalid.

Map **128**, in one embodiment, is used to map (i.e., translate) logical addresses to physical addresses within storage device **130**. Accordingly, as data becomes moved and invalidated, it may reside in different physical addresses on storage device **130** over time. Through the use of map **128**, however, an application may be able access a most recent set of data by specifying the same logical address (e.g., LBA) even though two or more versions of the data may reside in different physical addresses. Map **128** may be implemented using any suitable data structure. According, in one embodiment, map **128** is a binary-tree data structure. In others embodiments, map **128** may be an array, a linked list, a hash table, etc. In some embodiments, map **128** may be implemented using multiple data structures. Embodiments of map **128** are described in further detail below in conjunction with FIGS. **5** and **6**.

Applications **122**, in one embodiment, include program instructions that are executable by processor unit **110**. As will be described below, applications **122** may utilize various hardware of computing system such as processor unit **110**, RAM **120**, and storage device **130**. An operating system or hypervisor may allocate portions of storage device **130** and/or portions of RAM **120** to applications **122**.

Driver **126**, in one embodiment, is executable to permit applications **122** to interact with storage device **130**. Accordingly, driver **126** may receive requests to perform read and write operations at specified logical block addresses and may issue corresponding commands to controller **132** to implement those operations. In some embodiments, driver **126** manages garbage collection for storage device **130** to reclaim storage blocks with invalid data. As used herein, “reclaiming” a storage block or “reclamation” of a storage block refers to preparing the storage block for reuse (i.e., so that the storage block can store new data). In the case of flash media, reclamation may include copying valid data out of the storage block and erasing the block. In some embodiments, to facilitate performance of read and write operations, driver **126** also maps logical addresses (e.g., LBAs) to corresponding physical addresses (in other embodiments, mapping logical addresses to physical addresses may be performed elsewhere, such as at controller **132**). Accordingly, driver **126** may also manage map **128** including adding and removing translations from map **128** as data is manipulated on storage device **130**.

In some embodiments, driver **126** is configured to create and handle snapshots for data stored on storage device **130**.

Exemplary snapshot data is described below with reference to FIG. 2A. For example, driver 126 may provide a separate logical address range and epoch number for each snapshot. Data stored for each snapshot, in some embodiments, may include or indicate a delta of information written since a previous snapshot. Driver 126 may handle snapshots and epoch numbers for storage requests and translate logical addresses from applications 122 to physical addresses of storage device 130 based on the epoch numbers. The snapshot capability may or may not be transparent to applications 122. For example, in some embodiments, applications 122 may be configured to provide epoch numbers with particular storage requests. In other embodiments, applications 122 may provide storage requests without any knowledge of snapshots or epoch numbers.

In various embodiments, driver 126 presents a logical address space to applications 122. In one embodiment, the size of the logical address space may be equivalent to the size of the physical address space on storage device 130. For example, if storage device 130 has a 1.2 TB capacity addressable using a 32-bit physical address space, driver 126 may present a 32-bit logical address space to the operating system. In another embodiment, driver 126 presents a logical address space that is larger than the physical address space of storage device 130. In such an embodiment, applications 122 may be described as being “thinly provisioned” as they are given more resources (e.g., storage capacity) than actually exists—thus, applications 122 cannot collectively consume the entire logical address space (without adding additional capacity) as this would overload the storage capacity of storage device 130. Still further, in other embodiments, driver 126 may provide a logical address space that is significantly larger than the physical address space of a storage device such that the logical address space is a “sparse address space.” (For the purposes of this disclosure, a sparse address space is any logical address space that is at least 10% larger than the physical address space of a storage device.) For example, in one embodiment, driver 126 may present a 48-bit sparse address space relative to a 32-bit physical address space. In such an embodiment, a given application 122 may consume considerably less than its total allocated LBA range such that considerable unused portions of logical address space may exist between one application 122’s stored data and another application 122’s data. Driver 126 may determine the size of the logical address space to be presented based on any suitable criteria.

In various embodiments, allocating ranges of a larger logical address space may be advantageous because it reduces the possibility of collisions within the logical address space (e.g., applications 122 inadvertently accessing the same LBA). Allocated ranges may also be static, continuous, and non-overlapping to reduce the possibility of collisions. Still further, through the usage of map 128, driver 126 may reduce the possibility of collisions within the physical address space without relying on an operating system to prevent potential collisions.

Further, sparse addressing may allow for a larger number of snapshots relative to the actual available physical storage space. For example, less than the entire address range of a snapshot will typically be written during each epoch, so a physical address space may be smaller than the logical address space allocated to snapshots.

Turning now to FIG. 2A, a diagram illustrating exemplary snapshot data is shown. In the illustrated embodiment, snapshots 210A-C are each associated with a respective one of logical address ranges 215A-C, which are each used for write operations during a particular epoch. As used herein, the term

“epoch” refers to an interval between two events in time. Events may include snapshot events (e.g., creation, deletion, or merging of a snapshot), system events (e.g., power on, restart), particular points in time, etc. Epochs are often bounded by two consecutive snapshots. However, an “initial epoch” ends when a first snapshot occurs for a particular application, and may not be initially bounded by a snapshot (e.g., the initial epoch may start with an empty drive). Similarly, a “current epoch” occurs while performing storage requests after a previous snapshot, but is bounded by a present point in time rather than a snapshot until a subsequent snapshot occurs. Thus, storage requests to a storage range that implements snapshots occur during a particular epoch, which may be an initial epoch, a current epoch, or an epoch that is bounded by snapshots. Epochs may or may not be non-overlapping. For example, two epochs may overlap in the context where multiple child snapshots are created based on the same parent snapshot.

Further, as used herein, the term “snapshot” refers to the state of data associated with an application at a particular point in time. The term “snapshot data” refers to information that is stored to record and indicate a snapshot. Snapshot data may include metadata, mapping data, data itself, etc. in order to indicate the state of a storage device at a given time. For example, if snapshot data is generated for a particular application at time T1, the snapshot data can later be used at time T2 to recover data from time T1, even some of the data from T1 has been overwritten by the application at time T2. An initial snapshot may refer to an empty drive, which may be a logical address range with no valid data. Further, the phrase “creating a snapshot” refers to instantiating snapshot data. For example, creating a snapshot may include recording which data of the snapshot was written during a preceding epoch, pointing to data written during previous epochs, and/or allocating a new logical address range for data written during a current epoch that follows creation of the snapshot.

Driver 126, in various embodiments, is configured to assign a logical address range 215 to each epoch. In one embodiment, the logical address range for an initial epoch (not shown) is presented to an application. In this embodiment, the application may continue to use the logical address range of the initial epoch and may be unaware of other logical address ranges. Each logical address range may be located at a particular offset from the logical address range initially allocated to an application, and logical addresses from the application may be shifted based on epoch numbers in order to use the appropriate logical address range. In some embodiments, each logical address range 215 is the same size as the address space that is allocated to a particular application. An application may be unaware of the different logical address ranges 215, and may read and write from a single address range. Storage requests from the application may be directed to a particular logical address range 215 based on an epoch with which the logical address ranges 215 are associated.

In the illustrated example, data blocks 201 and 202 were written at blocks 6 and 14 of logical address range 215A during an epoch that preceded creation of snapshot N. Similarly, blocks 203 and 204 were written at blocks 5 and 20 of logical address range 115B during an epoch between creation of snapshot N and creation snapshot N+1, and blocks 205 and 206 were written at blocks 6 and 7 of logical address range 215C during the current epoch that began after snapshot N+1 was taken. Driver 126, in various embodiments, is configured to translate logical addresses in logical address ranges 215 to physical addresses on storage device 130. As described below with reference to FIG. 3A, physical blocks corresponding to

the logical addresses in logical address ranges **215** may be located at various different locations in a physical address space.

Thus in certain embodiments, for a given storage request, driver **126** may perform two translations: (1) a translation from an application's address to a logical address for the appropriate epoch and (2) a translation from the logical address for the epoch to a physical address.

In the illustrated embodiment, current data **220** is a conceptual overlay of data indicated by previous snapshot data. In this embodiment, snapshot data for a given snapshot indicates writes that occurred since the previous snapshot. To read current data, in one embodiment, driver **126** is configured to check a logical address range associated with a current epoch (logical address range **215C** in this example). If no translation exists for a block in the logical address range of the current epoch, (e.g., because the block has not been written during the current epoch), driver **126** is configured to examine the logical address ranges of one or more ancestor snapshots to find the data. This would occur when the block has not been written during the current epoch, for example. For example, when searching for the data at block **20**, driver **126** would stop when detecting block **204** in snapshot N+1. Note that in the example of FIG. 2A, block **206** "overwrites" block **201** at block **6** from the point of view of an application. Thus, block **206** would be read in response to a request for data from block **7** from current data and block **201** would be read in response to a request for data from block **7** from snapshot N+1. For a block at block **3**, for example, driver **126** may find the data in a snapshot previous to snapshot N, because the data does not have a translation for any of the snapshots shown.

In various embodiments, driver **126** is configured to assign an epoch number to each epoch. The epoch number may be used to identify the logical address range associated with an epoch. Driver **126** may handle shifting or adjusting addresses from an application to the appropriate logical address range based on these epoch numbers. Driver **126** may implement a counter to generate epoch numbers. The counter may wrap when it reaches a greatest representable value, making the number of snapshots available dependent on the number of bits used for the counter. In some embodiments, epoch number zero is associated with an initial snapshot which refers to an all-empty drive.

In some embodiments, driver **126** is configured to store an epoch number on storage device **130** at one or more locations that are associated with blocks written during the particular epoch. In one embodiment, as described below with reference to FIG. 3B, driver **126** is configured to store epoch numbers in metadata for each block (and/or for larger or smaller storage structures associated with each block). Each block may be a packet (described below with reference to FIG. 3B) or a larger data block. Driver **126** may also store epoch information in map **128** and/or other data structures for translating logical addresses to physical addresses.

Driver **126**, in one embodiment, is configured to create snapshots for applications at regular intervals. Alternatively or additionally, driver **126** may be configured to create snapshots based on requests by an application or operating system. The data included in a snapshot may be an entire logical address range available to an application, or some portion thereof. In some embodiments, an application may read from a particular epoch by providing an epoch number to driver **126**. In other embodiments, an application does not have access to epoch numbers and only sees data associated with overlay **240** unless a restoration of a previous snapshot occurs. In some embodiments, mapping of logical address

ranges **215** to physical addresses may persist for all snapshots that have not been deleted, allowing for easy rollback to a previous snapshot.

Driver **126** may delete snapshot data by issuing a TRIM command for the logical address range associated with the snapshot. Performing a TRIM command may include updating or removing nodes in map **128** associated with the logical address range and marking associated data stored on the storage device as invalid. A garbage collector may eventually re-use the physical locations associated with invalid data for new data. Thus, the data may be marked as invalid, but may not be deleted immediately. In some embodiments, driver **126** may wait to deleting data for a snapshot with multiple children until all blocks in its logical address range that have translations to physical addresses have been overwritten by the child snapshots, e.g., in order to avoid duplicating data.

Turning now to FIG. 2B, one embodiment of a data structure **280** for maintaining snapshot data is depicted. In the illustrated embodiment, nodes **281-288** in data structure **280** are each associated with a snapshot and contain references to respective parent and child snapshots. Each node may include information associated with the logical address range assigned to a snapshot, such as the bounds of the range, an offset associated with the range, or a pointer to information about the range, e.g., in map **128**. Each node may also include information about what blocks in the logical address range are valid. A logical address may be described as "valid" with respect to a snapshot when it is translatable to a corresponding physical address at which the data is stored, e.g., because the data was written during an epoch associated with the snapshot. In some embodiments, this validity corresponds to whether there is currently a physical address mapped to the logical address. In one embodiment, a particular logical address is implicitly invalid if there is no node in map **128** associated with the logical address.

Node **281**, in the illustrated embodiment, corresponds to an initial snapshot, which may represent an empty drive. Node **281**'s descendants include nodes for snapshot N, snapshot N+1 and the current epoch (CE). Node M **288**, in the illustrated embodiment, is another child node of node N+1 and is included to show that a snapshot may have multiple child snapshots. In the illustrated embodiment, each node includes references to both parent and child nodes (if existing), but in other embodiments, data structure **280** may include references in only one direction. In some embodiments, all snapshots for a particular data range share an oldest ancestor snapshot (e.g., initial snapshot **281** of FIG. 2B). Speaking generally, a "child snapshot" records data during an epoch immediately following creation of a "parent snapshot." Such parent/child relationships may be maintained using data structure **280** and/or some other structure.

Finding current data for a read operation may involve traversal of multiple nodes in the tree. For example, if a particular block is not valid in nodes **284** or **286**, a read may occur from snapshot N **282**. In some embodiments, driver **126** is configured to collapse the tree in order to improve efficiency. For example, driver **126** may merge parent snapshot data into its child by writing any current data from the parent snapshot forward and associating it with the child snapshot's epoch number. Merging snapshots may involve resolving conflicts between snapshots (e.g., when two snapshots include different data for the same block) using older or younger data in various configurations. In one embodiment, data from a younger snapshot is used when there is a conflict during merging. If a parent snapshot has multiple children, driver **126** may delete one or more of the children before performing such a merge, e.g., in order to prevent duplication of data.

11

Driver 126 may be configured to add a node to data structure 280 each time a snapshot is created and remove nodes from data structure 280 when snapshots are deleted. Driver 126 may balance data structure 280 in order to optimize access times.

Various optimizations may be implemented for performing read operations using data structure 280 and/or map 128. Speaking generally, driver 126 provides two keys (a logical address and an epoch number) to access a desired physical address on storage device 130. Thus, various optimizations used for databases that are accessed with multiple keys may be applied to driver 126.

In one embodiment, driver 126 is configured to maintain a unique forward map for the current epoch. This unique forward map may map logical address ranges to the epoch number in which they were written. For example, for the current epoch, this unique forward map may indicate, for each block in the logical address range, which snapshot holds valid data for the block. This implementation may improve access times by preventing traversals through multiple nodes to find current data.

In another embodiment, driver 126 may include reference data in map 128 that points to an epoch that contains valid data for a given block. For example, if data for a particular epoch does not contain a translation for a particular block, map 128 would include a reference to the epoch that contains the current data for the block. This implementation may reduce traversals of the map to two lookups at most but may increase the size of map 128.

In the illustrated embodiment, data structure 280 is a tree, but in other embodiments data structure 280 may be stored as an array, a linked list, a hash table, etc. Driver 126 may store data structure 280. In other embodiments, driver 126 may handle snapshot data without data structure 280, e.g., by using epoch numbers to access and update map 128. In some embodiments, driver 126 includes additional data structures (not shown) for optimizing storage requests associated with snapshots. Data structure 280 and map 128 may be separate, may be combined, and/or may reference each other, in various embodiments.

Turning now to FIG. 3A, an exemplary mapping of a logical address space 302 to a physical address space 304 is depicted. In one embodiment, logical address space 302 represents the organization of data as perceived by higher-level processes such as applications 122. In one embodiment, physical address space 304 represents the organization of data on the physical media.

Logical address space 302, in one embodiment, is divided into logical addresses corresponding to respective logical blocks 310A-310D (also referred to as sectors). In some embodiments, the logical addresses are LBAs (in other embodiments, the logical addresses may correspond to some other form of logical identifiers). In one embodiment, sectors/blocks 310 represent the smallest block of data associated with a given logical address. As but one example, a block 310 may be approximately 512 bytes in size (while logical erase blocks and logical pages discussed below may be approximately 40 MB and 8 kB, respectively).

Physical address space 304, in one embodiment, is divided into physical addresses corresponding to the arrangement of data on the physical recording media. As will be discussed in further detail with respect to FIG. 3B, in one embodiment, the content of logical blocks 310 may be stored as packets 360 within logical erase blocks 320. As discussed with respect to FIG. 3C, in various embodiments, physical address space 304 may be organized as a log structure, in which write operations may be performed at only one or more append points.

12

Turning now to FIG. 3B, a block diagram of storage blocks within storage device 130 is depicted. In the illustrated embodiment, storage device 130 is organized into logical erase blocks (LEBs) 320 that include multiple physical erase blocks (PEBs) 330, which are located in separate storage banks 134. A logical erase block 320 is further divided into multiple logical pages 340 (not to be confused with virtual memory pages) that, in turn, include multiple physical pages 350. Pages 350 include multiple packets 360, which may be grouped into ECC chunks 370.

As used herein, the term “erase block” refers broadly to a logical erase block or a physical erase block. In one embodiment, a physical erase block 330 represent the smallest storage block with a given bank 134 that can be erased at a given time (e.g., due to the wiring of cells on the die). In one embodiment, logical erase blocks 320 represent the smallest block erasable by controller 132 in response to receiving an erase command. In such an embodiment, when controller 132 receives an erase command specifying a particular logical erase block 320, controller 132 may erase each physical erase block 330 within the block 320 simultaneously. It is noted that physical erase blocks 330 within a given logical erase block 320 (e.g., blocks 330A and 330B) may be considered as contiguous in physical address space 304 even though they reside in separate banks 134. Thus, the term “contiguous” may be applicable not only to data stored within the same physical medium, but also to data stored within separate media.

In one embodiment, a physical page 350 represents the smallest storage block within a given bank 134 that can be written to at a given time. In one embodiment, a logical page 340 is the smallest writable storage block supported by controller 132. (In one embodiment, controller 132 may include a buffer configured to store up to a logical page worth of data; upon filling the buffer, controller 132 may write the contents of the buffer to a single logical page simultaneously.) In some instances, dividing a logical page 340 across multiple banks 134 may result in faster access times for a set of data when multiple banks 134 are accessed in parallel.

In one embodiment, a packet 360 represents the smallest storage block within a given bank 134 that can be read at a given time. In one embodiment, an ECC chunk 370 is the smallest storage block readable by controller 132. In some embodiments, packets 360 may be slightly larger than logical blocks 310 as they may include the contents of a logical block 310 (or multiple blocks 310 in some instances) as well as a packet header.

In some embodiments, driver 126 may associate metadata 390 with one or more of storage blocks 320-370. As used herein, the term “metadata” refers to system data usable to facilitate operation of solid-state storage device 130; metadata stands in contrast to, for example, data produced by an applications (i.e., “application data”) or forms of data that would be considered by an operating system as “user data.” For example, in one embodiment, a logical erase block 320 may include metadata specifying, without limitation, usage statistics (e.g., the number of program erase cycles performed on that block 320), health statistics (e.g., a value indicative of how often corrupted data has been read from that block 320), security or access control parameters, sequence information (e.g., a sequence indicator), a persistent metadata flag (e.g., indicating inclusion in an atomic storage operation), a transaction identifier, or the like. In one embodiment, the header within a packet 360 may include packet metadata such as one or more LBAs associated with the contained data, the packet size, linkages to other packets, error correction checksums, etc. In various embodiments, driver 126 may use this infor-

13

mation, along with other forms of metadata, to manage operation of storage device **130**. For example, driver **126** might use this information to facilitate performance of read and write operations, recover storage device **130** to a previous state (including, for example, reconstruction of various data structures used by driver and/or replaying a sequence of storage operations performed on storage device **130**), etc.

In one embodiment, driver **126** is configured to store an epoch number in metadata **390** for each packet that is part of a snapshot range. In this embodiment, the epoch number corresponds to an epoch during which the packet was written. In other embodiments, epoch numbers may be stored on storage device **130** and associated with other granularities. For example, an epoch number may be stored with each logical page visible to an application. As further examples, an epoch number may be written with each ECC chunk, each LEB, or any of various appropriate storage unit sizes. Metadata **390** may also include the logical address associated with each write operation. Driver **126** may use metadata **390** to determine when snapshot data should be merged or deleted, to reconstruct data structures (e.g., data structures that maintain relationships between snapshots), and restore storage device **130** to a previous state associated with a previous snapshot, e.g., when such data structures have been corrupted.

Turning now to FIG. 3C, a block diagram of log structure **380** within physical address space **304** is depicted. As shown, in various embodiments, data is stored sequentially at an append point **382** (also referred to as the “head”) that starts an initial logical page **340**. As additional data is stored, append point **382** advances to subsequent pages **340** in log structure **380**. Eventually, after storing enough data, the append point **382** reaches the “last” page **340** in storage device **130**, at which point the append point **382** wraps back to the initial page **340**. Thus, log structure **380** is depicted as a loop/cycle. As more data is stored, the number of available pages **340** (shown as unshaded pages **340**) decreases and the number of used pages **340** (shown as shaded pages **340**) increases. As discussed above, in order to reuse these pages **340** (i.e., make them available to receive further writes), in one embodiment, driver **126** performs erase operations on logical erase blocks **320**. In one embodiment, a tail **384** is maintained to identify the oldest page **340** still in use within structure **380** (pages other than the one located at the tail are considered to be younger than the tail). When the logical erase block **320** with the oldest page **340** is eventually erased, tail **384** is advanced forward to the next oldest page **340** in use at the end of log structure **380**.

In general, data that is modified less frequently than other data in storage device **130** will migrate towards tail **384** (such data may be described as having a “colder temperature” or simply as “cold data”). On the other hand, data that is modified more frequently (described as having a “hotter temperature” or as “hot” data) will typically be located closer to head **382**. Thus, valid data located in LEB **320A** is likely “colder” than data in LEB **320B**.

It is noted that, in other embodiments, storage device **130** may be organized in a non-log-structured format. As used herein, the term “strict log structure” refers to a structure in which write operations may be performed at only one append point, and are not allowed to “fill in” at locations behind the append point.

Turning now to FIG. 4, a block diagram illustrates one embodiment of driver **126**. As noted above, in one embodiment, driver **126** is executable to enable applications **122** to

14

interact with storage device **130**. In the illustrated embodiment, driver **126** includes a management layer **410** and an input/output (I/O) layer **420**.

Management layer **410**, in one embodiment, handles higher-level block-related operations for driver **126**. Accordingly, in various embodiments, management layer **410** tracks the mapping of logical addresses **402** to physical address **404**, and performs translations for addresses **402** received from higher-level processes such as those of applications **122** and/or an OS. In some embodiments, management layer **410** also performs garbage collection, e.g. using a groomer. In some embodiments, management layer **410** maintains various forms of metadata in one or more data structures located within RAM **120** such as forward map **414**, program erase (PE) statistics, health statistics, etc. (In other embodiments, data structures associated with driver **126** may be maintained elsewhere in system **100** such as within storage device **130**.) In one embodiment, driver **126** periodically stores copies of these data structures to storage device **130** so that they can be reconstructed in the event of a crash.

Forward map **414**, in one embodiment, is a forward mapping data structure usable to map a logical address space to a physical address. Forward map corresponds, in some embodiments, to map **128**. In some embodiments, forward map **414** may include metadata in addition to metadata used to facilitate mapping such as invalidity information. Although described as “forward” map, map **414** may also be used to perform a reverse mapping of a physical address to a logical address. As will be described in conjunction with FIGS. 5-6, in one embodiment, forward map **414** is implemented as an extended-range b-tree data structure. In other embodiments, forward map **414** may be implemented using other data structures such as arrays, hash tables, other forms of trees, etc. In some embodiments, forward map **414** may also be implemented as multiple data structures such as a tree with pointers into an array.

In the illustrated embodiment, forward map **414** is also configured to receive epoch numbers **406**. In some embodiments, epoch numbers **404** are encoded in logical addresses **402**. For example, the epoch numbers **406** may be included in the higher-order bits of logical addresses **402**. As another example, epoch numbers **404** may be encoded using a hash function or any of various appropriate algorithms for encoding an epoch number in a range of addresses. In other embodiments, epoch numbers **406** are provided separately from logical addresses **402**. Epoch numbers may be generated by driver **126** itself (e.g., using the current epoch) or may be received from an application (e.g., in order to access data from a previous snapshot). In some embodiments, forward map **414** is configured to generate physical addresses based on both logical addresses **402** and epoch numbers **406**.

I/O layer **420**, in one embodiment, handles lower-level interfacing operations with controller **132**. Accordingly, layer **420** may receive a write request or a read request and the physical address **404** for that request; layer **420** may then issue the appropriate commands to controller **132** to cause storage device **130** to fulfill that request. In some embodiments, I/O layer **420** may prepare data for DMA transactions and initialize a DMA controller to conduct the transactions.

Turning now to FIG. 5, a block diagram of map **128** is depicted. In illustrated embodiment, map **128** is an extended-range b-tree that includes multiple nodes **510A-C**. As shown, each node **510** includes a logical address range **520**, a physical address mapping **530**, one or more pointers **540**, additional metadata **550**, and an epoch number **406**.

In one embodiment, driver **126** is configured to allocate a new node in map **128** when a snapshot is created and the

15

current epoch begins. The node may include a logical address range that points to the entire logical address range associated with the current epoch. As data is written during the current epoch, the node may be divided up into multiple other nodes for smaller logical address portions within the range for the current epoch (see FIG. 6, for example).

Logical address range **520**, in one embodiment, is the range of logical addresses (e.g., LBAs) that are mapped using information within a given node **510**. Accordingly, logical address range **520A** specifies that physical address mapping **530A** pertains to LBAs 50-100, for example. If a logical address does not “hit” in a node **510** (i.e., does not fall within a range **520** of a node such as range **520A** in root node **510A**), then map **128** is traversed to examine ranges **520** in one or more leaf nodes such as nodes **510B** or **510C**. In one embodiment, map **128** includes a node **510** for each range of logical addresses that have been mapped to a corresponding range of physical addresses, but does not include nodes **510** corresponding to unmapped ranges. Thus, in such an embodiment, if a given LBA is unused, unallocated, and/or unwritten, a corresponding node **510** does not exist for that LBA in map **128**. On the other hand, if an LBA has been written to, map **128** includes a node **510** specifying range **520** that includes the LBA. As such, nodes **510** may be added and/or modified when data is written to storage device **130**. In such an embodiment, map **128** is also a sparse data structure, meaning that map **128** does not include mappings for an entire logical address space. Accordingly, in some embodiments, logical address space **302** may be significantly larger than physical address space **304**.

Physical address mapping **530**, in one embodiment, is the mapped physical addresses for a given range **520**. In one embodiment, a given physical address is a composite a bank identifier for a storage bank **134**, a PEB identifier for a PEB **330**, a physical page identifier for a page **350**, and a packet identifier for a packet **360**; however in other embodiments, a physical address may be organized differently (e.g., a composite of LEB, logical-page, and ECC-chunk identifiers). In one embodiment, physical address mapping **530** is specified as a range of physical addresses. In another embodiment, physical address mapping **530** is a base address that is combined with an offset determined from the logical address. In other embodiments, mapping **530** may be specified differently.

Pointers **540**, in one embodiment, identify leaf nodes **510** for a given node **510**. In some embodiments, map **128** is organized such that a left pointer identifies a node **510** that has a lower address range **520** than the present node **510** and a right pointer may identify a node **510** having a higher address range **520**. For example, if node **510A** corresponds to the logical address range 50-100, node **510B** may correspond to the range 0-50 and node **510C** may correspond to the range 100-150. In some embodiments, map **128** may also be periodically balanced to give it a logarithmic access time.

Metadata **550**, in one embodiment, is additional metadata that may not be used in mapping a logical address to physical address such as validity information and packet size. In one embodiment, validity information may identify whether particular locations (e.g., erase blocks, pages, or packets) store valid or invalid data. In some embodiments, metadata **550** may also include TRIM notes indicative of data that was invalidated in response to TRIM commands (in other embodiments, TRIM notes may be stored in a separate data structure within RAM **120**, or on storage device **130**). In some embodiments, storage device **130** may support variable packet sizes; in such an embodiment, metadata **550** may specify the size packets used for a given logical address range **520**. In some

16

embodiments, metadata **550** may also include other information such as age information, snapshot data, usage information (e.g., whether particular logical addresses are associated with hot or cold data), etc.

Epoch number **406**, in one embodiment is the epoch number associated with logical address range **520** (which may be all or a portion of a logical address range **215**, for example). In one embodiment, epoch number **560** is encoded in the upper bits of logical addresses. In another embodiment, epoch number **406** is included in metadata **550**. In still other embodiments, nodes in map **128** may not include epoch numbers and epoch numbers may be handled by other modules of driver **126**.

In some embodiments, driver **126** provides both an epoch number and a logical address to map **128**. As mentioned above, the epoch number may be encoded in the address, in some embodiments. Map **128**, in the illustrated embodiment, is configured to produce a physical address of storage device **130** in response to the received epoch number and logical address.

In one embodiment, driver **126** is configured to determine whether the snapshot data associated with the epoch number contains a translation to a physical address for the logical address before accessing map **128** (e.g., using data structure **280**). In this embodiment, driver **126** traverses map **128** to find the logical address and uses physical address mapping **530** to obtain the associated physical address. In another embodiment, driver **126** is configured to access map **128** before it has determined whether the snapshot data associated with the epoch number contains valid data for the logical address. In this embodiment, driver **126** may determine that the snapshot data associated with the epoch number does not contain valid data at the logical address based on metadata **440**. In this situation, driver **126** may be configured to re-traverse map **128** using a new logical address and/or new epoch number in order to access a node associated with a parent snapshot until a snapshot with valid data at the desired location is found.

Turning now to FIG. 6, another embodiment of map **128** is depicted. In the illustrated embodiment, epoch numbers are encoded in logical addresses (e.g., in the upper bits). In this embodiment, driver **126** may not include a separate data structure **280** for maintaining parent-child relationships between snapshots. In the illustrated embodiment, map **128** includes nodes **612-622** which are allocated among two snapshots **610**.

In the illustrated embodiment, nodes **612-622** are associated with snapshots based on bits of their address ranges (e.g., in higher-order bits of each address). Thus, a given snapshot may be associated multiple nodes in map **128** that encode the epoch number of the snapshot in their logical address range **520**.

Consider a situation in which snapshot N+1 **610B** is a child of snapshot N **610A**. Driver **126** accesses map **128** using a logical address that encodes an epoch number. If the node in the current snapshot N+1 associated with the logical address (e.g., node **618**) does not contain a valid entry, driver **126**, in one embodiment, is configured to apply an offset to the logical address. In this embodiment, driver **126** would apply the offset such that a new logical address would include snapshot N's epoch number and re-traverse map **128** in order to determine whether parent snapshot N included valid data for the new logical address.

Turning now to FIG. 7, a block diagram of driver **126** is depicted. In this embodiment, in contrast to the embodiments of FIGS. 4-6, driver **126** does not provide epoch numbers to a forward map. Driver **126** may be otherwise configured simi-

17

larly to the embodiment described above with reference to FIG. 4, with the addition of snapshot module 730.

Snapshot layer 730, in one embodiment, generates a new logical address 708 based on a logical address 402 and an epoch number 406. For example, snapshot layer 730 may be configured to apply offsets to logical addresses 402 based on epoch numbers 406 in order to generate new logical addresses 708. In one embodiment, layer 730 may use data structure 280 to translate epoch number 406 and logical address 402 to a corresponding logical address 708 usable to index into map 414. In other embodiments, snapshot layer 730 includes other data structures, such as a unique forward map for the current epoch that maps logical address ranges to the epoch number in which they were written. As discussed above, epoch numbers 406 may be generated by driver 126, an application 122, and/or some other module. In the illustrated embodiment, forward map 414 may not include epoch numbers 406.

Turning now to FIG. 8A, a method 800 is depicted. Method 800 is one embodiment of a method that may be performed by an apparatus such as computing system 100 or storage device 130. Accordingly, in one embodiment, the apparatus may execute program instructions of a driver such as driver 126 to perform method 800.

In step 810, a computing system maintains information that indicates the state of data associated with an application at a particular point in time. The computer system may maintain this information based on a periodic snapshot schedule or at the request of an application, for example. Maintaining the information may include assigning a logical address range to a current epoch that follows the particular point in time.

In step 820, the computing system assigns an epoch number to a current epoch. The current epoch may be a time interval between the particular point in time and a later point in time. The computing system may assign the epoch number using a counter and may assign epoch numbers sequentially. In one embodiment, the epoch number is encoded in one or more address bits (e.g., the higher-order bits) of a logical address range associated with the current epoch.

In step 830, the computing system writes a block of data to the storage device during the current epoch. In this embodiment, writing the block of data includes writing the epoch number to the storage device with the set of data. The computing system may write the epoch number in metadata included with a packet of data. The epoch number written to the storage device may be used to reconstruct snapshot information in the event of data corruption. The epoch number written to the storage device may also be used to determine when a snapshot should be deleted and/or merged into a child snapshot, for example.

Turning now to FIG. 8B, a flow diagram of a method 840 is depicted. Method 840 is one embodiment of a method that may be performed by an apparatus such as computing system 100 or storage device 130. Accordingly, in one embodiment, the apparatus may execute program instructions of a driver such as driver 126 to perform method 840.

In step 850, a storage device is configured such that the storage device has a logical address space that is larger than a physical address space of the storage device. This sparse addressing may allow computing system 100 to allocate more space for snapshots than actually exists on a storage device.

In step 860, a first write operation is performed during a first epoch using a first logical address range within the logical address space.

In step 870, a second write operation is performed during a second, later epoch using a second logical address range within the logical address space. The first and second write

18

operations may target the same location from the point of view or an application or may target different locations.

Turning now to FIG. 9, a block diagram of an apparatus 900 that includes modules is depicted. As used herein, the term “module” refers to circuitry configured to perform operations or a memory having program instructions stored therein that are executable by one or more processors to perform operations. Accordingly, a module may be implemented as a hardware circuit implemented in a variety of ways. The hardware circuit may include, for example, custom very-large-scale integration (VLSI) circuits or gate arrays, off-the-shelf semiconductors such as logic chips, transistors, or other discrete components. A module may also be implemented in programmable hardware devices such as field programmable gate arrays, programmable array logic, programmable logic devices, or the like. A module may also be any suitable form of non-transitory computer readable media storing program instructions executable to perform specified operations. Apparatus 900 may include modules to implement any of the functionality described herein. For example, in the illustrated embodiment, apparatus 900 includes an association module 910 and a storage module 920.

Association module 910, in one embodiment, is configured to associate a logical address range for a storage device with an epoch number, where the epoch number is associated with writes to the storage device during a particular epoch. The logical address range may be part of a logical address space in a sparse addressing configuration. In some embodiments, storage module 1120 may implement functionality described with respect to driver 126.

Storage module 920, in one embodiment, is configured to handle one or more storage operations associated with the epoch number using the logical address range. For example, in one embodiment, storage module 920 uses addresses in the logical address range for write operations and maps the addresses to physical locations on storage device 130. Storage module 920 may direct read operations to the logical address range based on determining that the read operations are associated with the epoch number. Storage module 920 may direct read operations to other logical address ranges associated with one or more ancestor snapshots in response to determining that a block within the logical address range targeted by the read operation is not valid in the current snapshot data. In some embodiments, storage module 920 may implement functionality described with respect to driver 126, storage device 130, or a combination thereof.

In some embodiments, association module 910 and/or storage module 920 are within a controller such as controller 132. In another embodiment, modules 910 and/or 920 may be located within a memory such as memory 120. In sum, the modules of apparatus 900 may be implemented in any suitable manner to perform functionality described herein. Apparatus 900 may also correspond to any suitable structure having the functionality of modules 910-920. In one embodiment, apparatus 900 is a computing system that includes (or is coupled to) a storage such as storage device 130. In another embodiment, apparatus 900 is a card including a controller (such as controller 132) and one or more storage elements (such as storage banks 134). In yet another embodiment, apparatus 900 is a computing system including a memory system that stores modules 910 and/or 920.

Turning now to FIG. 10A, a block diagram of an apparatus 1000 that includes a determination means 1010, a translation means 1020, and a storage means 1030 is depicted. Apparatus 1000 may correspond to any suitable structure having the functionality of determination means 1010, translation means 1020, and storage means 1030. For example, apparatus 1000

may be any suitable type of computing device such as a server, laptop, desktop, a mobile device, etc. In some embodiments, apparatus **1000** may include multiple computing devices working together. In some embodiments, apparatus **1000** is a card including a controller (such as controller **132**) and one or more storage elements (such as storage banks **134**).

In various embodiments, determination means **1010** may implement any of the functionality described herein with respect to driver **126**. Accordingly, in one embodiment, determination means **1010** determines a logical address for a write operation based on a current epoch number. For example, in one embodiment, determination means **1010** applies an offset to an address of the write operation based on the epoch number to determine the logical address. In one embodiment, determination means is configured to determine the logical address by encoding the epoch number in a logical address. For example, the epoch number may be encoded in the upper bits of the logical address. In some embodiments, determination means **1010** may also implement functionality other than that described in conjunction with driver **126**.

In various embodiments, translation means **1020** may implement any of the functionality described herein with respect to driver **126**. Accordingly, in one embodiment, translation means translates the logical address determined by determination means **1010** to a physical address or location on a storage device. For example, on one embodiment, translation means traverses map **128** in order to determine a physical address mapping **530** for the logical address. In some embodiments, determination means **1010** may also implement functionality other than that described in conjunction with driver **126**.

Determination means **1010** and translation means **1020** may correspond to any suitable structures. In one embodiment, determination means **1010** and/or translation means **1020** are hardware circuits configured to perform operations (e.g., controller **132**). The hardware circuits may include, for example, custom very-large-scale integration (VLSI) circuits or gate arrays, off-the-shelf semiconductors such as logic chips, transistors, or other discrete components. Means **1010** and/or **1020** may also be implemented in programmable hardware devices such as field programmable gate arrays, programmable array logic, programmable logic devices, or the like. In another embodiment, presentation means **1010** and/or **1020** includes a memory having program instructions stored therein (e.g., RAM **120**) that are executable by one or more processors (e.g., processor unit **110**) to implement an algorithm.

In one embodiment, determination means **1010** implements the algorithm discussed with respect to FIG. **10B**. In some embodiments, translation means **1020** implements the algorithm discussed with respect to FIG. **10C**. In some embodiments, determination means **1010** and/or translation means **1020** correspond to association module **910**. Accordingly, the phrases “means for determining a logical address for a write operation based a current epoch number” and “means for translating the logical address to a physical address” refer to any of the structures listed above as well as their corresponding equivalents.

In various embodiments, storage means **1030** may implement any of the functionality described herein with storage device **130**. Accordingly, in one embodiment, storage means **1030** is for storing the epoch number on the storage device. Storage means **1030** may store the epoch number at the same physical location as a target of the write operation or may store the epoch number at another location that is associated with the target of the write operation. Storage means **1030**

may correspond to any suitable structure such as those discussed above with respect to storage device **130** (e.g., one or more banks **134**, computing system **100**, storage system **200**, etc.). Accordingly, the phrase “means for storing the current epoch number on the storage device” refers to any of the structures listed above as well as their corresponding equivalents.

Turning now to FIG. **10B**, a flow diagram illustrating an algorithm **1030** is depicted. Algorithm **1030** is one embodiment of an algorithm implemented by determination means **1010**. In the illustrated embodiment, algorithm **1030** includes, at step **1032**, receiving a write operation and an epoch number. Algorithm **1030** further includes, at step **1034**, producing a logical address for the write operation based on the epoch number. The logical address may be produced by applying an offset to a logical address, by looking up the logical address in a data structure, and so on.

Turning now to FIG. **10C**, a flow diagram illustrating an algorithm **1050** is depicted. Algorithm **1050** is one embodiment of an algorithm implemented by translation means **1020**. In the illustrated embodiment, algorithm **1050** includes, at step **1052**, accessing a data structure using a logical address. For example, translation means **1020** may access map **128** and reading a physical address mapping. Algorithm **1050** further includes, at step **1054**, generating a physical address based on accessing the data structure. The physical address may be produced based on a log implementation.

Turning now to FIG. **11**, a block diagram of a storage system **1100** including computing system **100** is depicted. As discussed above, computing system **100** may include applications **122** that operate on data stored in storage device **130**. In the illustrated embodiment, computing system **100** executes a storage server application **1110** to enable client systems **1120A** and **1120B** to access and store data in storage device **1130** via network **1230**. For example, in one embodiment, storage system **1100** may be associated within an enterprise environment in which server application **1110** distributes enterprise data from storage device **1130** to clients **1120**. In some embodiments, clients **1120** may execute other server applications such as web servers, mail servers, virtual private network (VPN) servers, etc. to further distribute data to other computing systems. Accordingly, in some embodiments, storage server application **1110** may implement various network attached storage (NAS) protocols such as the file transfer protocol (FTP), network file system (NFS) protocol, server message block (SMB) protocol, Apple file protocol (AFP), etc. In some embodiments, computing system **100** may be one of several computing systems **100** configured to implement a storage area network (SAN).

Although specific embodiments have been described above, these embodiments are not intended to limit the scope of the present disclosure, even where only a single embodiment is described with respect to a particular feature. Examples of features provided in the disclosure are intended to be illustrative rather than restrictive unless stated otherwise. The above description is intended to cover such alternatives, modifications, and equivalents as would be apparent to a person skilled in the art having the benefit of this disclosure.

The scope of the present disclosure includes any feature or combination of features disclosed herein (either explicitly or implicitly), or any generalization thereof, whether or not it mitigates any or all of the problems addressed herein. Accordingly, new claims may be formulated during prosecution of this application (or an application claiming priority thereto) to any such combination of features. In particular, with ref-

21

erence to the appended claims, features from dependent claims may be combined with those of the independent claims and features from respective independent claims may be combined in any appropriate manner and not merely in the specific combinations enumerated in the appended claims. 5

What is claimed is:

1. A method, comprising:
 - maintaining, by a computing system, information that indicates the state of data associated with an application at a particular point in time; 10
 - assigning, by the computing system, an epoch number to a current epoch, wherein the current epoch is an interval between the particular point in time and a future point in time;
 - writing, by the computing system during the current epoch, a block of data to the storage device using a logical address, wherein the writing includes storing the epoch number in the same page as the block of data and wherein one or more bits of the logical address encode the epoch number. 15
2. The method of claim 1, further comprising:
 - associating, by the computing system, a logical address range with the epoch number;
 - translating, by the computing system, an address of a storage request from the application during the current epoch to a logical address within the logical address range; and 25
 - storing, by the computing system, the logical address with the block of data.
3. The method of claim 2, further comprising: 30
 - the computing system providing a logical address space for a storage device on which the block of data is stored, wherein the logical address space is larger than a physical address space of the storage device, wherein the logical address range is part of the logical address space. 35
4. The method of claim 1, wherein a size of the block of data is a smallest addressable unit for a storage device to which the block is written.
5. The method of claim 1, wherein the information includes snapshot data for a snapshot that shares a common ancestor snapshot with one or more other snapshots. 40
6. The method of claim 1, further comprising:
 - servicing a read request during the current epoch by providing data written to a storage device during a previous epoch. 45
7. The method of claim 1, further comprising:
 - performing a storage request that is associated with the epoch number by:
 - determining an entry in a forward mapping structure based on the epoch number; and 50
 - translating a logical address of the storage request to a physical address of a storage device using the entry.
8. The method of claim 1, further comprising:
 - receiving a storage request and an epoch number associated with the storage request; and 55
 - applying an offset to an address of the storage request based on the epoch number.
9. The method of claim 1, further comprising:
 - merging, after an end of the current epoch, the information with other information that indicates the state of data at the beginning of a later epoch, wherein the merging includes associating data written during the current epoch with an epoch number of the later epoch. 60
10. The method of claim 1, further comprising:
 - reconstructing a data structure that indicates parent-child relationships between one or more snapshots based on the epoch number stored with the block of data. 65

22

11. An apparatus, comprising:
 - an association module configured to associate a logical address range for a storage device with an epoch number, wherein the epoch number is associated with write operations to the storage device during a particular epoch; and
 - a storage module configured to service one or more storage operations associated with the epoch number using the logical address range, wherein one or more bits of logical addresses in the logical address range encode the epoch number of the particular epoch, wherein at least a portion of the association module and storage module comprise one or more of a hardware circuit and program instructions stored on one or more non-transitory computer readable media.
12. The apparatus of claim 11, wherein the storage module is further configured to:
 - perform a write operation that includes writing a block of data to the storage device including storing the epoch number on the storage device in the same erase block as the block of data.
13. The apparatus of claim 11, wherein the storage module is further configured to:
 - perform a read operation using a different logical address range associated with an epoch number of a previous epoch in response to determining that the read operation is for data written during the previous epoch; and
 - translate a logical address within the different address range to a physical address of the storage device to perform the read operation.
14. The apparatus of claim 11, wherein the storage module is further configured to:
 - perform a write operation to a logical address in the logical address range; and
 - indicate that the logical address is associated with valid data for the particular epoch.
15. The apparatus of claim 11, wherein the storage module is further configured to:
 - delete data associated with the particular epoch by issuing a trim command for the logical address range.
16. A non-transitory computer readable medium having program instructions stored thereon, wherein the program instructions are executable by a computing system to cause the computing system to perform operations comprising:
 - configuring a storage device such that the storage device has a logical address space that is larger than a physical address space of the storage device;
 - performing a first write operation during a first epoch using a first logical address range within the logical address space, wherein one or more bits of logical addresses in the first logical address range encode a first epoch number; and
 - performing a second write operation during a second, later epoch using a second logical address range within the logical address space, wherein one or more bits of logical addresses in the second logical address range encode a second epoch number.
17. The non-transitory computer-readable storage medium of claim 16,
 - wherein the performing the first write operation during the first epoch includes storing an epoch number associated with the first epoch and a first logical address on the storage device at a location associated with the write operation, wherein the first logical address is within the first logical address range.
18. The non-transitory computer-readable storage medium of claim 16, wherein the operations further comprise:

applying an offset to an address of a storage request based on an epoch number provided with the storage request, wherein applying the offset provides a new logical address.

19. The non-transitory computer-readable storage medium of claim 16, wherein the operations further comprise: receiving a read request during the second epoch; determining that a logical address for the read request in the second logical address range does not have a valid translation to a physical address; and in response to the determining, servicing the read request using data from a physical location mapped to a logical address in the first address range.

20. The non-transitory computer-readable storage medium of claim 16, wherein the operations further comprise: handling storage requests to the storage device using a strict log structure.

21. An apparatus, comprising:

a first means for determining a logical address for a write operation based a current epoch number, wherein one or more bits of the logical address encode the current epoch number; a second means for translating the logical address to a physical address on a storage device; and a third means for storing the current epoch number on the storage device, wherein storing the current epoch number comprises storing the current epoch number at the same physical location as a target of the write operation.

* * * * *